

# optimization

November 29, 2022

## 1 Write a script that implement the GD algorithm, with the following structure:

Input:  $f$ : the function  $f(x)$  we want to optimize. It is supposed to be a Python function, not an array.  $\text{grad\_}f$ : the gradient of  $f(x)$ . It is supposed to be a Python function, not an array.  $x_0$ : an  $n$ -dimensional array which represents the initial iterate.  $k_{\max}$ : an integer. The maximum possible number of iterations (to avoid infinite loops)  $\text{tol}_f$ : small float. The relative tolerance of the algorithm. Convergence happens if  $\|\text{grad\_}f(x_k)\|_2 < \text{tol}_f \|\text{grad\_}f(x_0)\|_2$   $\text{tol}_x$ : small float. The tolerance in the input domain. Convergence happens if  $\|x_k - x_{k-1}\|_2 < \text{tol}_x$ . Pay attention to the first iterate. Output:  $x$ : an array that contains the value of  $x_k$  FOR EACH iterate  $x_k$  (not only the latter).  $k$ : an integer. The number of iteration needed to converge.  $k < k_{\max}$ .  $f\_val$ : an array that contains the value of  $f(x_k)$  FOR EACH iterate  $x_k$ .  $\text{grads}$ : an array that contains the value of  $\text{grad\_}f(x_k)$  FOR EACH iterate  $x_k$ .  $\text{err}$ : an array the contains the value of  $\|\text{grad\_}f(x_k)\|_2$  FOR EACH iterate  $x_k$ .

```
[71]: import matplotlib.pyplot as plt
import numpy
import numpy as np
```

```
[72]: def plot(x, errf, x_true=False, back=False, alpha=0.2):
    k = len(x)
    title = f"{'x*=' + str(x_true) if x_true else ''} x_c={np.round(x[-1], 2)} N. of iteration: {k}, backtracking: {'yes' if back else 'no'} {'alpha: ' + str(alpha) if not back else ''}"
    plt.title(title)
    plt.plot(errf)
    legend = ["error"]
    if x_true:
        x_errors = np.zeros(x.shape)
        for i, x_k in enumerate(x):
            x_errors[i] = np.linalg.norm(x_k - x_true)
        plt.plot(x_errors)
        legend.append("X error")
    #plt.subplot(2, 2, 2)
    #plt.plot(points, grads)
    #plt.subplot(2, 2, 3)
    #plt.plot(points, err)
```

```

plt.legend(legend)
plt.show()

def backtracking(f, grad_f, x):
    """
    This function is a simple implementation of the backtracking algorithm for
    the GD (Gradient Descent) method.

    f: function. The function that we want to optimize.
    grad_f: function. The gradient of f(x).
    x: ndarray. The actual iterate x_k.
    """
    alpha = 1
    c = 0.8
    tau = 0.25

    while f(x - alpha * grad_f(x)) > f(x) - c * alpha * np.linalg.
↪norm(grad_f(x), 2) ** 2:
        alpha = tau * alpha

        if alpha < 1e-3:
            break
    return alpha

def GD(f, grad_f, x0, tolf, tolx, kmax, alpha=0.2, back=False):

    x0 = np.array(x0)
    shape = (kmax, *x0.shape)
    # output
    x = np.zeros(shape)
    f_val = np.zeros(shape)
    grads = np.zeros(shape)
    err = np.zeros(shape)

    x_tol = tolx
    f_tol = tolf
    x_old = x0
    k = 0

    while k < kmax and x_tol >= tolx and f_tol >= tolf:
        if back:
            alpha = backtracking(f, grad_f, x_old)
            x_k = x_old - alpha * np.array(grad_f(x_old))
            x_tol = np.linalg.norm(x_k-x_old)

```

```

    f_tol = np.linalg.norm(f(x_k))

    # Update arrays
    x[k] = x_k
    f_val[k] = f(x_k)
    grads[k] = grad_f(x_k)
    err[k] = np.linalg.norm(grads[k])
    x_old = x_k
    k = k+1

    return x[:k], f_val[:k], grads[:k], err[:k]

```

## 1.1 Test the algorithm above on the following functions:

```

[73]: tolf = 1e-4
      tolx = 1e-4
      kmax = 100
      alphas = [0.1, 0.01]

```

### 1.1.1 Function 1

$$f(x_1, x_2) = (x_1 - 3)^2 + (x_2 - 1)^2$$

for which the true optimum is  $x^* = (3, 1)^T$

```

[74]: def f1(x):
      x1, x2 = x
      return (x1 - 3)**2 + (x2 - 1)**2

      def grad_f1(x):
          x1, x2 = x
          return np.array((2*(x1-3), 2*(x2-1)))

      x0 = (2, 2)
      x_true1 = (3, 1)

      def test_function(f, grad_f, x0, kmax, x_true=False, f5=False):
          for alpha in alphas:
              x, f_val, grads, err = GD(f, grad_f, x0, tolf, tolx, kmax, alpha)
              plot(x, err, x_true, alpha=alpha)
              if f5:
                  x_ = np.linspace(-3, 3, 1000)
                  plt.plot(x_, f(x_))
                  plt.plot(x, f_val, "bo")

```

```


plt.show()
x, f_val, grads, err = GD(f, grad_f, x0, tolf, tolx, kmax, back=True)
plot(x, err, x_true)

if f5:
    x_ = np.linspace(-3, 3, 1000)
    plt.plot(x_, f(x_))
    plt.plot(x, f_val, "bo")
    plt.show()


test_function(f1, grad_f1, x0, kmax, x_true1)

```


$x^* = (3, 1)$   $x_c = [2.99 \ 1.01]$  N. of iteration: 23, backtracking: no alpha: 0.1



$x^* = (3, 1)$   $x_c = [2.87 \ 1.13]$  N. of iteration: 100, backtracking: no alpha: 0.01



$x^* = (3, 1)$   $x_c = [2.99 \ 1.01]$  N. of iteration: 38, backtracking: no alpha: 0.2




## 1.2 Function 2


$$f(x_1, x_2) = 10(x_1 - 1)^2 + (x_2 - 2)^2$$

```
[75]: def f2(x):  
      x1, x2 = x  
      return 10*(x1 - 1)**2 + (x2 - 2)**2  
      def grad_f2(x):  
          x1, x2 = x  
          return np.array((20*(x1-1), 2*(x2-2)))  
  
      x0 = (2, 2)  
      x_true2 = (1,2)  
  
      test_function(f2, grad_f2, x0, kmax, x_true2)
```


x\* = (1, 2) x\_c = [2. 2.] N. of iteration: 100, backtracking: no alpha: 0.1



$x^* = (1, 2)$   $x_c = [1. \ 2.]$  N. of iteration: 26, backtracking: no alpha: 0.01



$x^* = (1, 2)$   $x_c = [1. \ 2.]$  N. of iteration: 16, backtracking: no alpha: 0.2



### 1.3 Function 3

$$f(x) = \frac{1}{2} \|Ax - b\|_2^2$$

```
[76]: def f3(x):
    x = np.array(x)
    x = np.reshape(x, (1, len(x)))
    n, m = x.shape
    x_true = np.ones((1, n))
    v = np.linspace(0, 1, n, endpoint=True)
    A = numpy.vander(v)
    b = A @ x_true
    return 1/2 * np.linalg.norm(A @ x - b, 2)**2

def grad_f3(x):
    n = len(x)
    v = np.linspace(0, 1, n)
    A = np.vander(v)
    x_true = np.ones(n).T
    b = A @ x_true
```




```
return np.array(A.T@(A@x-b))
```

```
N = [5, 10, 15]
```


```
for n in N:  
    x0 = [3 for i in range(n)]  
    print("N = ", n)  
    test_function(f3, grad_f3, x0, kmax)
```

```
N = 5
```


x\_c=[1.06 1.02 0.98 0.95 1.02] N. of iteration: 100, backtracking: no alpha: 0.1



$x_c=[1.25 \ 1.17 \ 1.07 \ 0.91 \ 0.83]$  N. of iteration: 100, backtracking: no alpha: 0.01

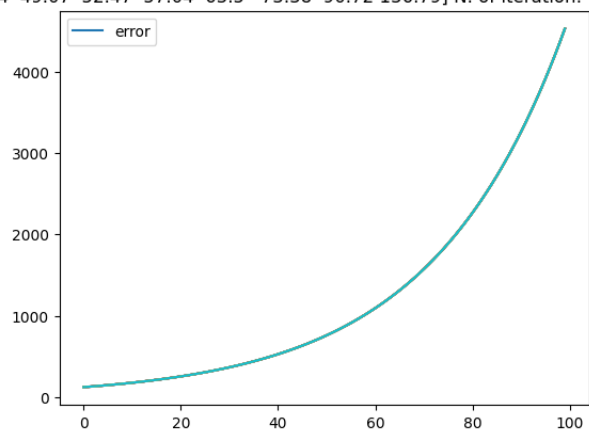


$x_c=[2.18 \ 2.12 \ 2.03 \ 1.87 \ 1.43]$  N. of iteration: 100, backtracking: no alpha: 0.2




N = 10


x\_c=[ 42.68 44.36 46.44 49.07 52.47 57.04 63.5 73.38 90.72 136.79] N. of iteration: 100, backtracking: no alpha: 0.1



x\_c=[1.17 1.14 1.1 1.06 1.01 0.96 0.9 0.85 0.85 1.09] N. of iteration: 100, backtracking: no alpha: 0.01




x\_c=[1.81 1.77 1.72 1.67 1.6 1.52 1.4 1.25 1.01 0.54] N. of iteration: 100, backtracking: no alpha: 0.2




N = 15


x\_c=[5.99208298e+34 6.15224507e+34 6.33555016e+34 6.54693822e+34  
6.79290087e+34 7.08216634e+34 7.42679457e+34 7.84400901e+34  
8.35943550e+34 9.01324025e+34 9.87285441e+34 1.10627513e+35  
1.28473569e+35 1.59313784e+35 2.37801057e+35] N. of iteration: 100, backtracking: no alpha: 0.1



$x_c=[1.14 \ 1.13 \ 1.1 \ 1.08 \ 1.06 \ 1.03 \ 1.01 \ 0.98 \ 0.95 \ 0.92 \ 0.9 \ 0.89 \ 0.9 \ 0.95$   
 $1.08]$  N. of iteration: 100, backtracking: no alpha: 0.01



$x_c=[1.63 \ 1.6 \ 1.57 \ 1.54 \ 1.5 \ 1.46 \ 1.41 \ 1.36 \ 1.29 \ 1.21 \ 1.12 \ 1. \ 0.86 \ 0.67$   
 $0.42]$  N. of iteration: 100, backtracking: no alpha: 0.2



## 1.4 Function 4

$$f(x) = \frac{1}{2} \|Ax - b\|_2^2 + \frac{\lambda}{2} \|x\|_2^2$$

```
[77]: def f4_builder(lmb):
    def f4(x):
        x = np.array(x)
        x = np.reshape(x, (1, len(x)))
        n, m = x.shape
        x_true = np.ones((1, n))
        v = np.linspace(0, 1, n, endpoint=True)
        A = numpy.vander(v)
        b = A @ x_true
        return 1/2 * np.linalg.norm(A @ x - b, 2)**2 + lmb/2 * np.linalg.
        ↪norm(x)**2
    return f4


def grad_f4_builder(lmb):
    def grad_f4(x):
        n = len(x)
        v = np.linspace(0,1,n)
        A = np.vander(v)
        x_true = np.ones(n).T
        b = A @ x_true
        return np.array(A.T@(A@x-b)) + lmb*x
    return grad_f4

n = 5
lmbs = np.linspace(0, 1, 3)
x0 = [0 for i in range(n)]


for lmb in lmbs:
    print("Lambda: ", lmb)
    f4 = f4_builder(lmb)
    grad_f4 = grad_f4_builder(lmb)
    test_function(f4, grad_f4, x0, kmax)
```

Lambda: 0.0


$x_c=[0.97 \ 0.99 \ 1.01 \ 1.03 \ 0.99]$  N. of iteration: 100, backtracking: no alpha: 0.1



$x_c=[0.87 \ 0.91 \ 0.97 \ 1.04 \ 1.08]$  N. of iteration: 100, backtracking: no alpha: 0.01




$x_c=[0.41 \ 0.44 \ 0.49 \ 0.57 \ 0.79]$  N. of iteration: 100, backtracking: no alpha: 0.2




Lambda: 0.5

$x_c=[0.84 \ 0.87 \ 0.92 \ 0.98 \ 1.02]$  N. of iteration: 41, backtracking: no alpha: 0.1






$x_c=[0.8 \ 0.84 \ 0.89 \ 0.98 \ 1.07]$  N. of iteration: 100, backtracking: no alpha: 0.01




$x_c=[0.4 \ 0.43 \ 0.48 \ 0.56 \ 0.77]$  N. of iteration: 100, backtracking: no alpha: 0.2




Lambda: 1.0


$x_c=[0.76 \ 0.79 \ 0.84 \ 0.92 \ 1.02]$  N. of iteration: 26, backtracking: no alpha: 0.1



$x_c=[0.74 \ 0.78 \ 0.83 \ 0.92 \ 1.05]$  N. of iteration: 100, backtracking: no alpha: 0.01




x\_c=[0.39 0.42 0.47 0.54 0.76] N. of iteration: 100, backtracking: no alpha: 0.2

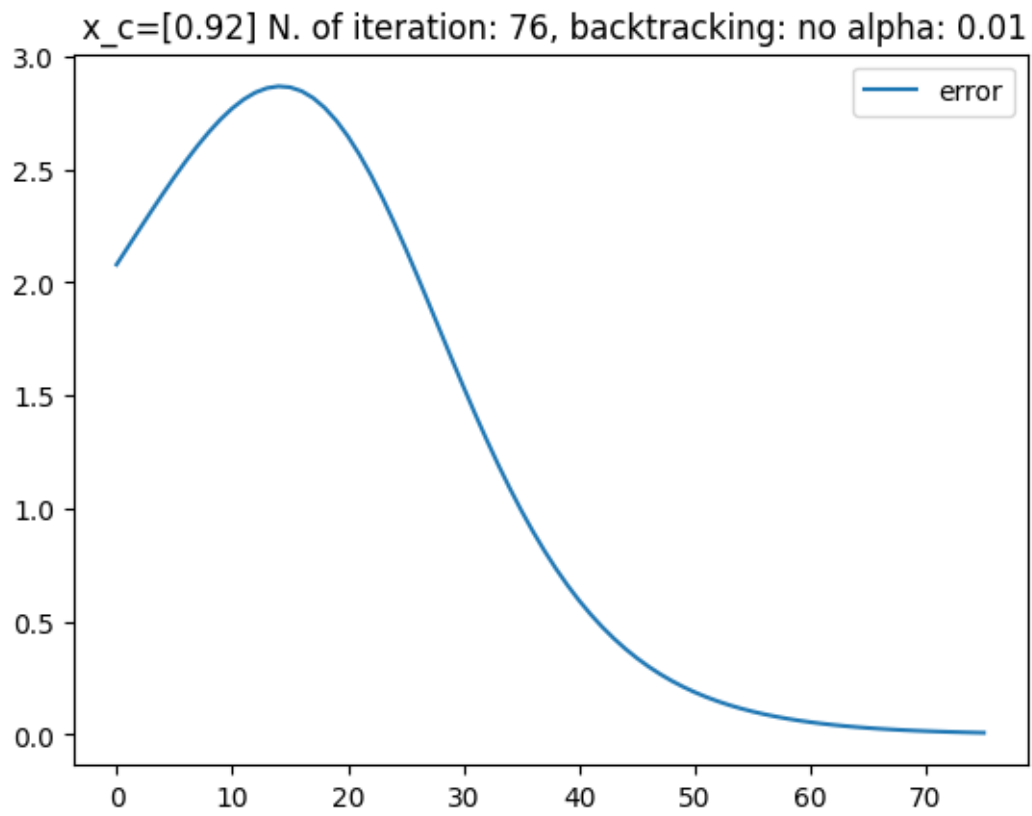



#### 1.4.1 Function 5

$$f(x) = x^4 + x^3 - 2x^2 - 2x$$

```
[78]: def f5(x):  
        return np.power(x,4) + np.power(x, 3) - 2*np.power(x,2) - 2*x  
  
    def grad_f5(x):  
        return np.array(4*np.power(x, 3) + 3*np.power(x,2) - 4*x - 2)  
  
    N = np.arange(5, 20, 5)  
  
    x0 = [0.]  
    test_function(f5, grad_f5, x0, kmax)
```









Only for the non-convex function defined in 5, plot it in the interval  $[-3, 3]$  and test the convergence point of GD with different values of  $x_0$  and different step-sizes. Observe when the convergence point is the global minimum and when it stops on a local minimum or maximum.

```
[79]: x_5 = np.linspace(-3, 3, 1000)
plt.plot(x_5, f5(x_5))
plt.show()


starting_points = [-2, 0, 2]


for x0 in starting_points:
    x0 = np.array([x0])
    test_function(f5, grad_f5, x0, kmax, f5=True)
```






















































$x_c=[0.92]$  N. of iteration: 41, backtracking: no alpha: 0.2







Hard (optional): For the functions 1 and 2, plot the contour around the minimum and the path defined by the iterations (following the example seen during the lesson). See `plt.contour` to do that.

```
[80]: def contour(f, grad_f, x0, x_true, radius, tol_x, tol_f, kmax):
    x11, x12 = x_true


    xv = np.linspace(x11 - radius, x11 + radius, 30)
    yv = np.linspace(x12 - radius, x12 + radius, 30)
    xx, yy = np.meshgrid(xv, yv)


    x, f_val, grads, err = GD(f, grad_f, x0, tol_f, tol_x, kmax, back=True)
    zz = f1((xx, yy))

    plt.figure(figsize=(10,10))
    plt.contour(xx, yy, zz)
    plt.plot(grads[:, 0], grads[:, 1], 'o-')
    plt.title("Counter plot")
    plt.grid()
    plt.show()

contour(f1, grad_f1, (0, 0), x_true1, 10, tol_x, tol_f, kmax)
```

```
contour(f2, grad_f2, (0, 0), x_true2, 20, tol_x, tol_f, kmax)
```





## 2 Optimization via Stochastic Gradient DescentInput:

$l$ : the function  $l(w; D)$  we want to optimize. It is supposed to be a Python function, not an array.  
 $\text{grad}_l$ : the gradient of  $l(w; D)$ . It is supposed to be a Python function, not an array.  $w_0$ : an  $n$ -dimensional array which represents the initial iterate. By default, it should be randomly sampled.  
 $\text{data}$ : a tuple  $(x, y)$  that contains the two arrays  $x$  and  $y$ , where  $x$  is the input data,  $y$  is the output data.  
 $\text{batch\_size}$ : an integer. The dimension of each batch. Should be a divisor of the number of data.  
 $n\_epochs$ : an integer. The number of epochs you want to repeat the iterations. Output:  
 $w$ : an array that contains the value of  $w_k$  FOR EACH iterate  $w_k$  (not only the latter).  $f\_val$ : an array that contains the value of  $l(w_k; D)$  FOR EACH iterate  $w_k$  ONLY after each epoch.  
 $\text{grads}$ : an array that contains the value of  $\text{grad}_l(w_k; D)$  FOR EACH iterate  $w_k$  ONLY after each epoch.  
 $\text{err}$ : an array the contains the value of  $\|\text{grad}_l(w_k; D)\|_2$  FOR EACH iterate  $w_k$

ONLY after each epoch.

```
[81]: # Import the data MNIST

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import scipy as sp

dataset = pd.read_csv("data.csv")
dataset = np.array(dataset)

def split(X, Y, Ntrain):
    d, N = X.shape

    idx = np.arange(N)
    np.random.shuffle(idx)

    train_idx = idx[:Ntrain]
    test_idx = idx[Ntrain:]

    Xtrain = X[:, train_idx]
    Ytrain = Y[train_idx]

    Xtest = X[:, test_idx]
    Ytest = Y[test_idx]

    return (Xtrain, Xtest, Ytrain, Ytest)

[82]: def create_dataset(dataset, digits=[3,6], Ntrain=4600):
    digits = [3, 6]

    y = dataset[:, 0]
    x = dataset[:, 1:].T

    Y = y#y.reshape((len(y), 1))
    X = np.concatenate((np.ones((1, len(y))), x), axis=0)

    I1 = (Y == digits[0]) # (Y[:, 0] == digits[0])
    I2 = (Y == digits[1]) # (Y[:, 0] == digits[1])
    X1 = X[:, I1]
    X2 = X[:, I2]
    Y1 = np.zeros((len(Y[I1]), ))
    Y2 = np.ones((len(Y[I2]), ))
    #Y1 = Y[I1]
    #Y2 = Y[I2]
```

```

X = np.concatenate((X1, X2), axis=1)
Y = np.concatenate((Y1, Y2))

d, N = X.shape

#Ntrain = 4600#int(N/3*2)

x_train, x_test, y_train, y_test = split(X, Y, Ntrain)
return x_train, x_test, y_train, y_test

```

```

[83]: x_train, x_test, y_train, y_test = create_dataset(dataset)
print(x_train.shape, x_test.shape, y_train.shape)

```

```

(785, 4600) (785, 3888) (4600,)

```

```

[84]: def sigmoid(x):
        return 1 / (1 + np.exp(-x))

def f(x, w):
    x_cup = x
    return sigmoid(x_cup.T @ w)

def grad_f(x):
    return np.exp(-x) / (np.exp(-x) + 1) ** 2

def l(w, x, y):
    return np.mean(np.linalg.norm(f(x, w)-y)**2)

def loss_grad(w, x, y):
    d, n = x.shape

    y = np.array(y)

    sum = 0
    for i in range(n):
        z = f(x[:, i], w)
        sum += z * (1 - z) * x[:, i].T * (z - y[i])

    return sum / n

def grad_l(w, x, y):
    x_cup = x

```

```

    return np.mean(sigmoid(x_cup.T @ w)*(1-sigmoid(x_cup.T @ w)) * x_cup.T *
↪(f(x, w) - y))

```

```
[85]: n_epochs = 50
```

```
[86]: batch_size = 15
```

```

def batch(x, y, batch_size):
    n = x.shape[1]
    idx = np.arange(n)
    np.random.shuffle(idx)
    n_batches = n // batch_size
    for i in range(n_batches):
        batch_index = idx[i*batch_size:(i+1)*batch_size]
        yield x[:, batch_index], y[batch_index]

```

```
def SGD(l, grad_l, w0, data, batch_size, n_epochs=50):
```

```

    shape = (n_epochs, *w0.shape)
    x_, y_ = data
    d, n = x_.shape
    w = np.zeros(shape)
    f_val = np.zeros((n_epochs, 1))
    grads = np.zeros((n_epochs, 1, w0.shape[0]))
    err = np.zeros((n_epochs, 1))

```

```
    #print("W0 shape: ", w0.shape)
```

```
    alpha = 1e-2
```

```
    w_old = w0
```

```
    w_k = w_old
```

```

    for epoch in range(n_epochs):
        batch_iterator = batch(x_, y_, batch_size)
        for x, y in batch_iterator:
            grad = grad_l(w_old, x, y)
            w_k = w_old - alpha * grad
            w_old = w_k
        w[epoch] = w_k
        f_val[epoch] = l(w_k, x_, y_)
        grads[epoch] = grad_l(w_k, x_, y_)
        err[epoch] = np.linalg.norm(grads[epoch], 2)
    return w, f_val, grads, err

```

```
data = (x_train, y_train)
```

```

sigma = 1e-3
d, N = x_train.shape
w0 = np.random.normal(0, sigma, (d, ))
w_sgd, f_val__sgd, grads_sgd, err_sgd = SGD(
    1, loss_grad, w0, data, batch_size, n_epochs)

```

/tmp/ipykernel\_11714/3753579620.py:2: RuntimeWarning: overflow encountered in exp

```

    return 1 / (1 + np.exp(-x))

```

```

[87]: def predict(w, X, treshhold=0.5):
        y_pred = f(X, w)
        y_copy = np.copy(y_pred)
        y_pred[y_copy < treshhold] = 0
        y_pred[y_copy >= treshhold] = 1
        return y_pred

def accuracy(y_hat, y):
    return round(np.mean(y_hat == y), 10)

```

```

[88]: w_star_sgd = w_sgd[-1]
        y_hat = predict(w_star_sgd, x_train)
        print("Accuracy: ", accuracy(y_hat, y_train))
        x_plot = np.arange(n_epochs)
        plt.plot(x_plot, err_sgd)
        plt.show()

```


/tmp/ipykernel\_11714/3753579620.py:2: RuntimeWarning: overflow encountered in exp

```

    return 1 / (1 + np.exp(-x))

```

Accuracy: 0.9936956522



```
[89]: def GD2(f, grad_f, D, w0, tolf=1e-9, tolx=1e-9, kmax=50, alpha=0.1, back=False):

    x_cup, y = D
    shape = (kmax, *w0.shape)

    w = np.zeros(shape)
    f_val = np.zeros((kmax, 1))
    grads = np.zeros((kmax, 1, w0.shape[0]))
    err = np.zeros((kmax, 1))
    # output

    #x_cup = np.concatenate((np.ones((1, N)), X), axis=0)

    x_tol = tolx
    f_tol = tolf
    w_old = w0
    k = 0

    while k < kmax and x_tol >= tolx and f_tol >= tolf:
        if back:
            alpha = backtracking(f, grad_f, w_old)
```



```

w_k = w_old - alpha * grad_f(w_old, x_cup, y)
x_tol = np.linalg.norm(w_k-w_old)
f_tol = np.linalg.norm(f(w_k, x_cup, y))

# Update arrays
w[k] = w_k
f_val[k] = f(w_k, x_cup, y)
grads[k] = grad_f(w_k, x_cup, y)
err[k] = np.linalg.norm(grads[k])
w_old = w_k
k = k+1

return w[:k], f_val[:k], grads[:k], err[:k]

```

```

[97]: sigma = 1e-3
w0 = np.random.normal(0, sigma, (d, ))
w_gd, f_val_gd, grads_gd, err_gd = GD2(1, loss_grad, (x_train, y_train), w0)

```

```


/tmp/ipykernel_11714/3753579620.py:2: RuntimeWarning: overflow encountered in
exp
    return 1 / (1 + np.exp(-x))

```

```

[98]: w_star_gd = w_gd[-1]
x_plot = np.arange(len(w_gd))
plt.plot(x_plot, err_gd)
plt.show()
y_hat = predict(w_star_gd, x_train)
print("Accuracy: ", accuracy(y_hat, y_train))

```



Accuracy: 0.9531102733

```
/tmp/ipykernel_11714/3753579620.py:2: RuntimeWarning: overflow encountered in
exp
return 1 / (1 + np.exp(-x))
```

```
[92]: x_train, x_test, y_train, y_test = create_dataset(dataset, digits=[1, 7],
↳ Ntrain=dataset.shape[0]//2)
```

```
[103]: data = (x_train, y_train)
sigma = 1e-3
d, N = x_train.shape
w0 = np.random.normal(0, sigma, (d, ))
w_sgd, f_val_sgd, grads_sgd, err_sgd = SGD(
    1, loss_grad, w0, data, batch_size, n_epochs)

sigma = 1e-3
w0 = np.random.normal(0, sigma, (d, ))
w_gd, f_val_gd, grads_gd, err_gd = GD2(1, loss_grad, (x_train, y_train), w0)
```


```
/tmp/ipykernel_11714/3753579620.py:2: RuntimeWarning: overflow encountered in
exp
return 1 / (1 + np.exp(-x))
```

```
[104]: w_star_sgd = w_sgd[-1]
y_hat = predict(w_star_sgd, x_train)
print("Accuracy: ", accuracy(y_hat, y_train))
x_plot = np.arange(n_epochs)
plt.plot(x_plot, err_sgd)
plt.show()
```


/tmp/ipykernel\_11714/3753579620.py:2: RuntimeWarning: overflow encountered in exp

```
    return 1 / (1 + np.exp(-x))
```

Accuracy: 0.9917530631



```
[105]: w_star_gd = w_gd[-1]
x_plot = np.arange(len(w_gd))
plt.plot(x_plot, err_gd)
plt.show()
y_hat = predict(w_star_gd, x_train)
print("Accuracy: ", accuracy(y_hat, y_train))
```



Accuracy: 0.8225730443

```
/tmp/ipykernel_11714/3753579620.py:2: RuntimeWarning: overflow encountered in  
exp  
    return 1 / (1 + np.exp(-x))
```